

PROJECT (3)

August 1, 2024

1 Project Requirements

- Final dataset should include 3 changes (deleting, inserting, adding, etc.)
- Pandas, matplotlib, and scikit packages should be part of the solution.
- Able to display results based on scenarios (For example, if I have different output, how the result will change. It could be a graph or display of results in row format)
- Minimum of 5 charts for insights. Use Python libraries to draw graphs.
- Total ten insights from the dataset

1.0.1 The Final Report should include the following.

- Description of project
- Overview of the dataset
- Exploratory analysis
- Details of insights will be identified from the dataset.
- Python programs for all 10 insights
- Minimum of 5 charts

1.1 1. Abstract

Airbnb is a prominent rental property service catering to vacationers, winter avoiders, getaway seekers, and travelers worldwide. This project delves into the insights behind Airbnb listings in the United States. We will analyze the top cities with the most rental properties, evaluate the pricing trends across different states, and develop a predictive model for estimating rental prices.

```
[1]: #Import the required packages:

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

#import data using pandas

df = pd.read_csv("airbnb_listings_usa.csv", low_memory=False)
```

1.2 Overview of Data and Exploratory Data Analysis (EDA)

To begin we need to get a better understanding of the data. We will look at things such as number of rows, columns, descriptive variables to help describe the data, and even visualizations to get a better understanding of how the data is formatted in the file. Below is the column names and descriptions from the dataset. To view the dataset please see (<https://www.kaggle.com/datasets/tamle507/airbnb-listings-usa>)

Column Name	Description
id	Unique identifier for the listing
name	Name of the listing
host_id	Unique identifier for the host
host_name	Name of the host
neighbourhood_group	Broad geographical area of the listing (e.g., borough or city district)
neighbourhood	Specific neighbourhood where the listing is located
latitude	Latitude coordinate of the listing
longitude	Longitude coordinate of the listing
room_type	Type of room offered (e.g., entire home/apt, private room, shared room)
price	Price per night for the listing (in local currency)
minimum_nights	Minimum number of nights required to book the listing
number_of_reviews	Total number of reviews the listing has received
last_review	Date of the last review received
reviews_per_month	Average number of reviews per month
calculated_host_listings_count	Total number of listings the host has
availability_365	Number of days the listing is available for booking in a year
number_of_reviews_ltm	Number of reviews in the last twelve months
license	License information for the listing (if applicable)
state	State where the listing is located
city	City where the listing is located

[2]: *#head, this will show us the top 5 rows in the dataset. Giving us a good idea of the layout of columns and what they may include*

```
df.head()
```

```
[2]: Unnamed: 0      id      name \
0      0  183319  Panoramic Ocean View Venice Beach
1      1   109  Amazing bright elegant condo park front *UPGRA...
2      2  51307  Spanish Bungalow Guest House LA CA. 30 plus ni...
3      3 184314  Boho Chic Flat..Steps to Beach!
4      4  51498  Guest House With Its Own Entrance/Exit and Hot...

      host_id host_name neighbourhood_group  neighbourhood  latitude \
0    867995  Barbara X  City of Los Angeles      Venice  33.99211
1      521    Paolo    Other Cities      Culver City  33.98301
```

2	235568	David	City of Los Angeles	Atwater Village	34.12206
3	884031	Ashley	City of Los Angeles	Venice	33.97487
4	236758	Bay	City of Los Angeles	Mar Vista	34.00389

	longitude	room_type	...	minimum_nights	number_of_reviews	\
0	-118.47600	Entire home/apt	...	30	3	
1	-118.38607	Entire home/apt	...	30	2	
2	-118.26783	Entire home/apt	...	30	138	
3	-118.46312	Entire home/apt	...	30	30	
4	-118.44126	Entire home/apt	...	3	378	

	last_review	reviews_per_month	calculated_host_listings_count	\
0	2019-02-25	0.02	2	
1	2016-05-15	0.01	1	
2	2020-12-13	0.98	2	
3	2017-12-24	0.22	1	
4	2022-08-21	2.60	1	

	availability_365	number_of_reviews_ltm	license	state	city
0	0	0	NaN	CA	Los Angeles
1	139	0	NaN	CA	Los Angeles
2	224	0	NaN	CA	Los Angeles
3	0	0	NaN	CA	Los Angeles
4	348	41	HSR19-001336	CA	Los Angeles

[5 rows x 21 columns]

```
[3]: #this command allows us to quickly see all column names in the dataset
df.columns
```

```
[3]: Index(['Unnamed: 0', 'id', 'name', 'host_id', 'host_name',
        'neighbourhood_group', 'neighbourhood', 'latitude', 'longitude',
        'room_type', 'price', 'minimum_nights', 'number_of_reviews',
        'last_review', 'reviews_per_month', 'calculated_host_listings_count',
        'availability_365', 'number_of_reviews_ltm', 'license', 'state',
        'city'],
        dtype='object')
```

```
[4]: #This will show the shape in (rows, column) order
df.shape
```

```
[4]: (325858, 21)
```

```
[5]: #use of dtypes shows the type of data in each column, either int64 or integer/
      ↪ numeric, object/string, or float64 which is decimal value
df.dtypes
```

```
[5]: Unnamed: 0          int64
      id                int64
      name              object
      host_id           int64
      host_name         object
      neighbourhood_group object
      neighbourhood     object
      latitude          float64
      longitude         float64
      room_type         object
      price             int64
      minimum_nights    int64
      number_of_reviews int64
      last_review       object
      reviews_per_month float64
      calculated_host_listings_count int64
      availability_365   int64
      number_of_reviews_ltm int64
      license           object
      state             object
      city              object
      dtype: object
```

```
[6]: #Describe the numeric columns in the dataframe
      df.describe()
```

```
[6]:
```

	Unnamed: 0	id	host_id	latitude \
count	325858.000000	3.258580e+05	3.258580e+05	325858.000000
mean	162928.500000	1.541106e+17	1.446528e+08	34.676058
std	94067.246346	2.736013e+17	1.449951e+08	6.213029
min	0.000000	1.090000e+02	2.300000e+01	18.920250
25%	81464.250000	2.394619e+07	2.311836e+07	32.775202
50%	162928.500000	4.511097e+07	9.320864e+07	34.102360
75%	244392.750000	5.420558e+07	2.388338e+08	39.948015
max	325857.000000	7.251653e+17	4.810023e+08	47.748000

	longitude	price	minimum_nights	number_of_reviews \
count	325858.000000	325858.000000	325858.000000	325858.000000
mean	-106.354815	284.915304	13.430175	39.457850
std	24.176674	835.569711	28.783033	75.724832
min	-159.714620	0.000000	1.000000	0.000000
25%	-118.410259	97.000000	2.000000	1.000000
50%	-117.131590	159.000000	3.000000	9.000000
75%	-82.549423	275.000000	30.000000	42.000000
max	-70.913250	100000.000000	1250.000000	2600.000000

	reviews_per_month	calculated_host_listings_count	availability_365 \
--	-------------------	--------------------------------	--------------------

count	263166.00000	325858.000000	325858.000000
mean	1.69220	27.108900	182.361099
std	2.01294	79.983052	134.095115
min	0.01000	1.000000	0.000000
25%	0.33000	1.000000	55.000000
50%	1.01000	2.000000	174.000000
75%	2.49000	12.000000	322.000000
max	190.48000	660.000000	365.000000

	number_of_reviews_ltm
count	325858.000000
mean	11.889817
std	20.672830
min	0.000000
25%	0.000000
50%	3.000000
75%	16.000000
max	1284.000000

```
[7]: #describe the categorical columns in the dataframe
df.describe(include = 'O')
```

	name	host_name	neighbourhood_group	\
count	325839	324714	155047	
unique	264019	33120	34	
top	Boutique Hostel	Blueground	City of Los Angeles	
freq	152	4876	39002	

	neighbourhood	room_type	last_review	license	state	\
count	325146	325858	263166	86635	325858	
unique	1575	4	2965	48064	19	
top	Unincorporated Areas	Entire home/apt	2022-09-05	Exempt	CA	
freq	13400	243098	13190	6603	127329	

	city
count	325858
unique	31
top	Los Angeles
freq	91630

2 Cleaning

Now that we have a general understanding of the layout of the dataset, we can dive deeper into the data. This will consist of handling missing values, duplicates, drop columns, add columns, visualizations.

```
[8]: #delete unneeded columns and missing values
df = df.drop(['Unnamed: 0', 'id', 'host_id', 'license', 'neighbourhood_group'],
            ↪axis = 1)
df.isna().sum()
```

```
[8]: name                19
     host_name           1144
     neighbourhood       712
     latitude            0
     longitude           0
     room_type           0
     price               0
     minimum_nights      0
     number_of_reviews   0
     last_review         62692
     reviews_per_month   62692
     calculated_host_listings_count  0
     availability_365     0
     number_of_reviews_ltm  0
     state               0
     city                0
     dtype: int64
```

```
[9]: #drop rows where missing values in name, host_name, last review,
     ↪reviews_last_month, neighbourhood as these are probalby no longer active
     ↪airbnb

df.dropna(subset=['name', 'host_name', 'last_review', 'reviews_per_month',
            ↪'neighbourhood'], inplace = True)

df.isna().sum()
```

```
[9]: name                0
     host_name           0
     neighbourhood       0
     latitude            0
     longitude           0
     room_type           0
     price               0
     minimum_nights      0
     number_of_reviews   0
     last_review         0
     reviews_per_month   0
     calculated_host_listings_count  0
     availability_365     0
     number_of_reviews_ltm  0
     state               0
```

```
city                                0
dtype: int64
```

```
[10]: df.shape
```

```
[10]: (262040, 16)
```

```
[11]: #identify outliers in the data

import numpy as np

# Define numeric features for outlier detection
numeric_features = ['price', 'minimum_nights', 'number_of_reviews',
                    'reviews_per_month', 'calculated_host_listings_count',
                    'availability_365']

# Function to detect outliers using IQR within each city
def detect_outliers_iqr(df, features):
    outliers = pd.DataFrame()
    non_outliers = pd.DataFrame()
    for city in df['city'].unique():
        city_df = df[df['city'] == city]
        Q1 = city_df[features].quantile(0.25)
        Q3 = city_df[features].quantile(0.75)
        IQR = Q3 - Q1
        city_outliers = ((city_df[features] < (Q1 - 1.5 * IQR)) |
            (city_df[features] > (Q3 + 1.5 * IQR))).any(axis=1)
        outliers = pd.concat([outliers, city_df[city_outliers]])
        non_outliers = pd.concat([non_outliers, city_df[~city_outliers]])
    return non_outliers, outliers

# Identify and remove outliers
df_no_outliers, df_outliers = detect_outliers_iqr(df, numeric_features)
print(f'Number of records before removing outliers: {len(df)}')
print(f'Number of outliers identified and removed: {df_outliers.shape[0]}')

print(f'Number of records after removing outliers: {df_no_outliers.shape[0]}')

import matplotlib.pyplot as plt
import seaborn as sns

# Plot box plots for each city
plt.figure(figsize=(15, 10))
```

```

sns.boxplot(x='city', y='price', data=df)
plt.xticks(rotation=45)
plt.title('Box Plot of Prices for Each City')
plt.show()

# Scatter plot to identify outliers in price vs. other features within each city
plt.figure(figsize=(15, 10))
sns.scatterplot(x='minimum_nights', y='price', hue='city', data=df,
               palette='Set1')
plt.title('Scatter Plot of Price vs Minimum Nights by City')
plt.show()


plt.figure(figsize=(15, 10))
sns.scatterplot(x='number_of_reviews', y='price', hue='city', data=df,
               palette='Set1')
plt.title('Scatter Plot of Price vs Number of Reviews by City')
plt.show()


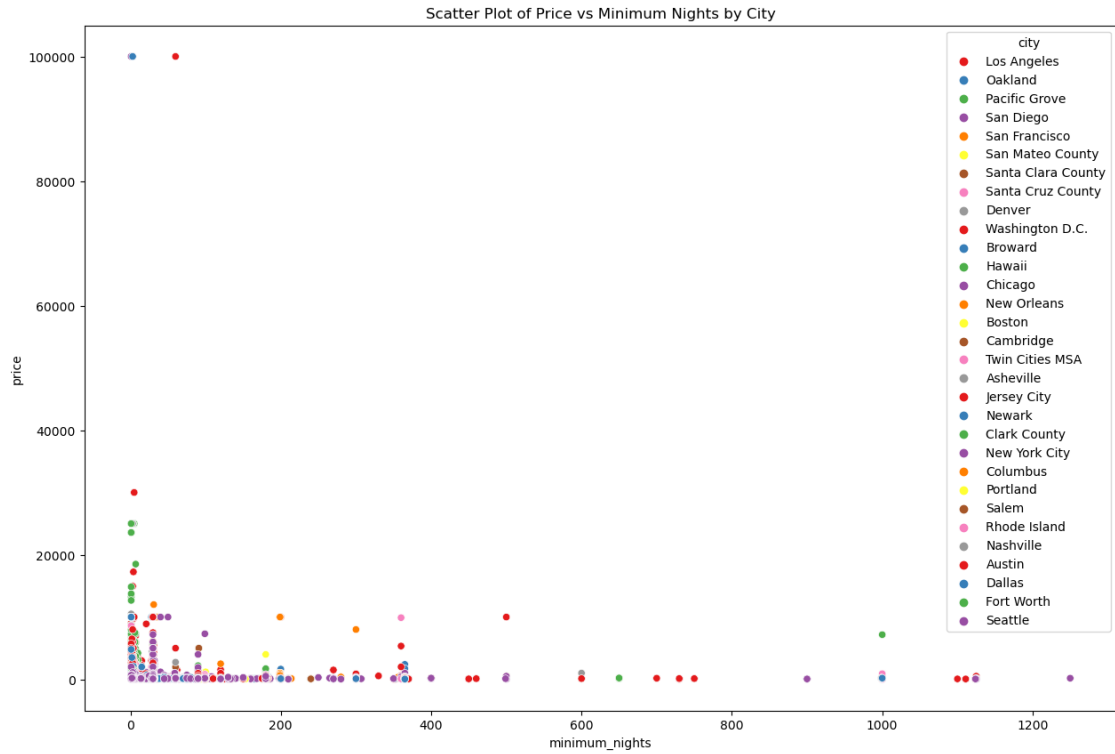
```

Number of records before removing outliers: 262040

Number of outliers identified and removed: 96708

Number of records after removing outliers: 165332

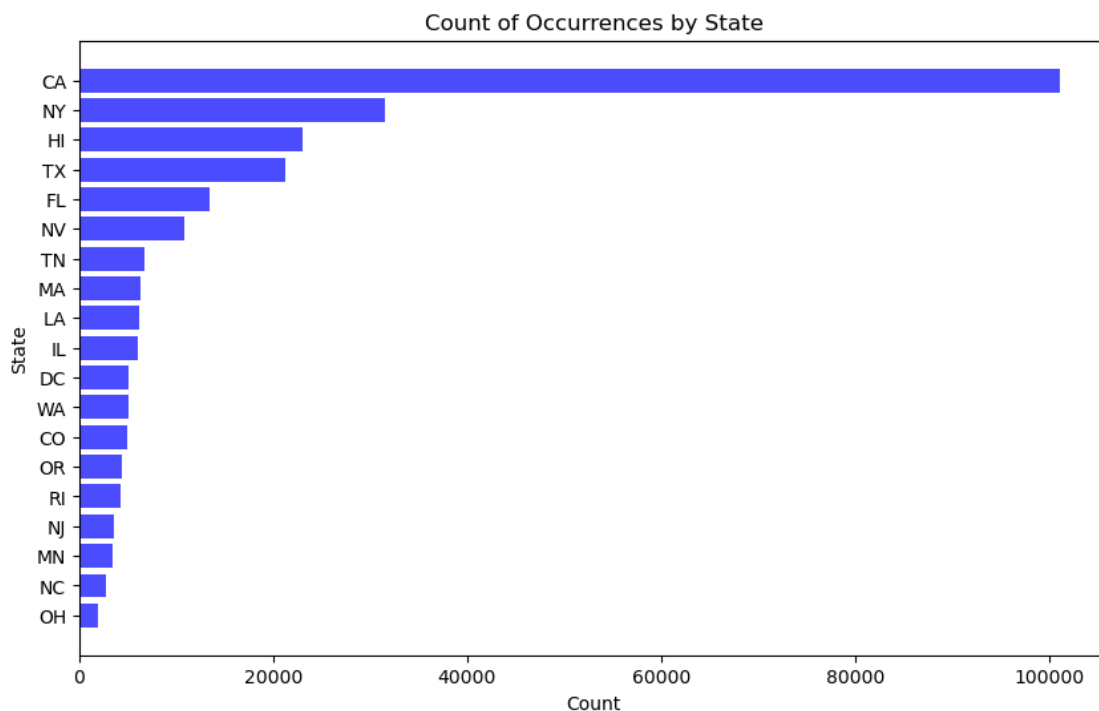




2.1 Findings

```
[12]: #figure showing distribution by state
state_counts = df.groupby('state').size().reset_index(name='count').
    ↪sort_values('count')

plt.figure(figsize=(10, 6))
plt.barh(state_counts['state'], state_counts['count'], color='blue', alpha=0.7)
plt.xlabel('Count')
plt.ylabel('State')
plt.title('Count of Occurrences by State')
plt.show()
```




As you can see from the graphs above, majority of the listing on Airbnb in this dataset are located in California. The top 5 states with listings are California, New York, Hawaii, and Texas. California outnumbers all other states by almost 3 times.

```
[13]: state = df.groupby('state')['price'].agg(['count', 'mean', 'median', 'max',
    ↪ 'min']).sort_values(by=['count', 'mean', 'median'], ascending =
    ↪ [False, False, False])
state = state.sort_values('median', ascending = False)

# #plot average price per state
plt.figure(figsize = (10,6))
```

```
plt.bar(state.index,state['median'],color = 'lightgreen')
plt.title('Average Price per State')
plt.xlabel('State')
plt.ylabel('Average Price ($)')
plt.show()
```



From the following you can see Rhode Island has the highest median price of listings. Following Rhode Island, the top 5 rounds out with Hawaii , Tennessee, Massachusents, and Nevada. As a potential investor considering Texas, you can see that the median price is on the lower side, and the state is top 5 in listings. I would consider this a good investment area because Texas is vast and large with many major areas, spreading out the competition. However we want to get a better understanding of what city is best to inveset into.

```
[14]: # Ensure we are working with a fresh copy
texas = df[df['state'] == 'TX'].copy()

# Create a popularity score using .loc
texas.loc[:, 'popularity_score'] = texas['number_of_reviews'] +
    ↪(texas['reviews_per_month'] * 30) - (texas['availability_365'] / 10)

# Sort properties by popularity score
most_popular_properties = texas.sort_values(by='popularity_score',
    ↪ascending=False).reset_index()
```

```

# Select the top 10 most popular properties
top_10_most_popular = most_popular_properties.head(10)

print(top_10_most_popular[['name', 'state', 'city']])

# Plot the top 10 most popular properties without special characters
plt.figure(figsize=(12, 6))
plt.barh(top_10_most_popular['name'], top_10_most_popular['popularity_score'],
         color='skyblue')
plt.xlabel('Popularity Score')
plt.title('Top 10 Most Popular Airbnb Properties in Texas')
plt.gca().invert_yaxis()
plt.show()

# Ensure we are working with a fresh copy
california = df[df['state'] == 'CA'].copy()

# Create a popularity score using .loc
california.loc[:, 'popularity_score'] = california['number_of_reviews'] +
    (california['reviews_per_month'] * 30) - (california['availability_365'] /
    10)

# Sort properties by popularity score
most_popular_properties = california.sort_values(by='popularity_score',
         ascending=False).reset_index()

# Select the top 10 most popular properties
top_10_most_popular = most_popular_properties.head(15)

print(top_10_most_popular[['name', 'state', 'city']])

# Plot the top 10 most popular properties without special characters
plt.figure(figsize=(12, 6))
plt.barh(top_10_most_popular['name'], top_10_most_popular['popularity_score'],
         color='skyblue')
plt.xlabel('Popularity Score')
plt.title('Top 10 Most Popular Airbnb Properties in California')
plt.gca().invert_yaxis()
plt.show()

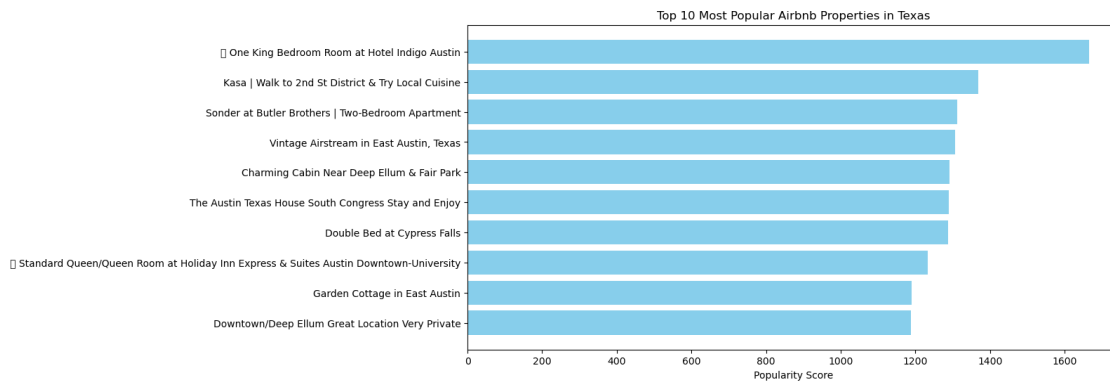
```

	name	state	city
0	One King Bedroom Room at Hotel Indigo Austin	TX	Austin
1	Kasa Walk to 2nd St District & Try Local Cui...	TX	Austin
2	Sonder at Butler Brothers Two-Bedroom Apartment	TX	Dallas

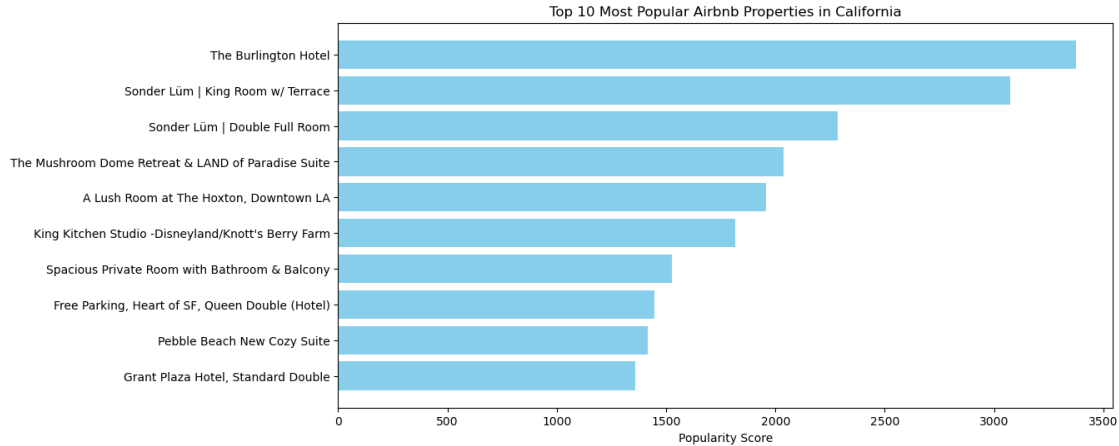
3	Vintage Airstream in East Austin, Texas	TX	Austin
4	Charming Cabin Near Deep Ellum & Fair Park	TX	Dallas
5	The Austin Texas House South Congress Stay and...	TX	Austin
6	Double Bed at Cypress Falls	TX	Austin
7	Standard Queen/Queen Room at Holiday Inn Exp...	TX	Austin
8	Garden Cottage in East Austin	TX	Austin
9	Downtown/Deep Ellum Great Location Very Private	TX	Dallas

/Users/brianhonea/anaconda3/lib/python3.11/site-packages/IPython/core/pylabtools.py:152: UserWarning: Glyph 10024 (\N{SPARKLES}) missing from current font.

```
fig.canvas.print_figure(bytes_io, **kw)
```



	name	state	city
0	The Burlington Hotel	CA	Los Angeles
1	The Burlington Hotel	CA	Los Angeles
2	Sonder Lüm King Room w/ Terrace	CA	Los Angeles
3	Sonder Lüm King Room w/ Terrace	CA	Los Angeles
4	Sonder Lüm Double Full Room	CA	Los Angeles
5	Sonder Lüm Double Full Room	CA	Los Angeles
6	The Mushroom Dome Retreat & LAND of Paradise S...	CA	Santa Cruz County
7	A Lush Room at The Hoxton, Downtown LA	CA	Los Angeles
8	A Lush Room at The Hoxton, Downtown LA	CA	Los Angeles
9	King Kitchen Studio -Disneyland/Knott's Berry ...	CA	Los Angeles
10	King Kitchen Studio -Disneyland/Knott's Berry ...	CA	Los Angeles
11	Spacious Private Room with Bathroom & Balcony	CA	San Francisco
12	Free Parking, Heart of SF, Queen Double (Hotel)	CA	San Francisco
13	Pebble Beach New Cozy Suite	CA	Pacific Grove
14	Grant Plaza Hotel, Standard Double	CA	San Francisco



2.1.1 Findings on the Most Popular Airbnb Properties in Texas

Based on our analysis of the Airbnb dataset, we created a popularity score for properties in Texas by considering the number of reviews, reviews per month, and availability. This score was calculated as follows:

$$\text{popularity_score} = \text{number_of_reviews} + (\text{reviews_per_month} * 30) - (\text{availability_365} / 10)$$

After sorting the properties by their popularity scores, we identified the top 10 most popular Airbnb properties in Texas. Here are the key findings:

1. **Top Properties:** The most popular properties are located in various cities across Texas, highlighting the widespread interest in Airbnb accommodations throughout the state.
2. **High Review Count:** These properties have a high number of reviews, indicating that they are frequently booked and reviewed by guests.
3. **Active Engagement:** Properties with higher reviews per month are consistently attracting guests, suggesting active engagement and positive experiences that encourage reviews.
4. **Low Availability:** Many of the top properties have lower availability throughout the year, implying they are frequently booked and in high demand.

This analysis provides valuable insights into the factors contributing to the popularity of Airbnb properties. Hosts and property managers can leverage this information to enhance their listings, improve guest experiences, and increase bookings by focusing on generating positive reviews and maintaining high occupancy rates.

```
[15]: # Group by city and calculate mean and median prices
texas_cities = texas.groupby('city')['price'].agg(['mean', 'median']).
      ↪reset_index()

# Define the positions for the bars
x = np.arange(len(texas_cities['city']))
width = 0.35
```

```

# Create the plot
plt.figure(figsize=(14, 8))
plt.bar(x - width/2, texas_cities['mean'], width, label='Mean Price')
plt.bar(x + width/2, texas_cities['median'], width, label='Median Price')

# Add labels, title, and legend
plt.xlabel('City')
plt.ylabel('Price')
plt.title('Mean and Median Prices of Airbnb Listings by City in Texas')
plt.xticks(ticks=x, labels=texas_cities['city'], rotation=45, ha='right')
plt.legend()

# Show the plot
plt.tight_layout()
plt.show()

# Group by city and calculate mean and median prices
california_cities = california.groupby('city')['price'].agg(['mean', 'median']).
    ↪reset_index()

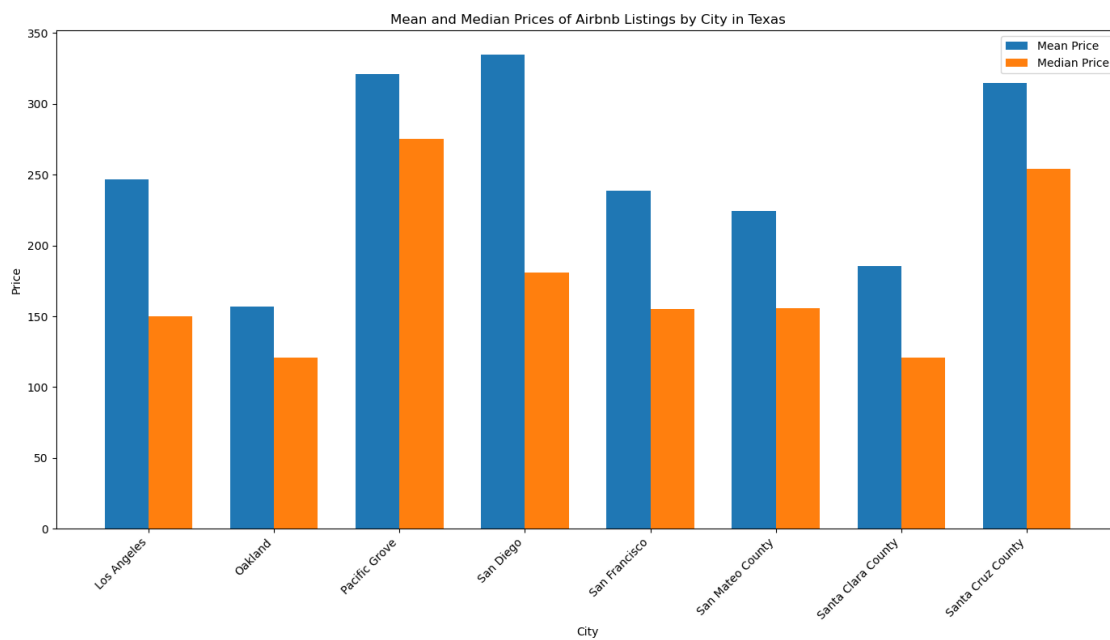
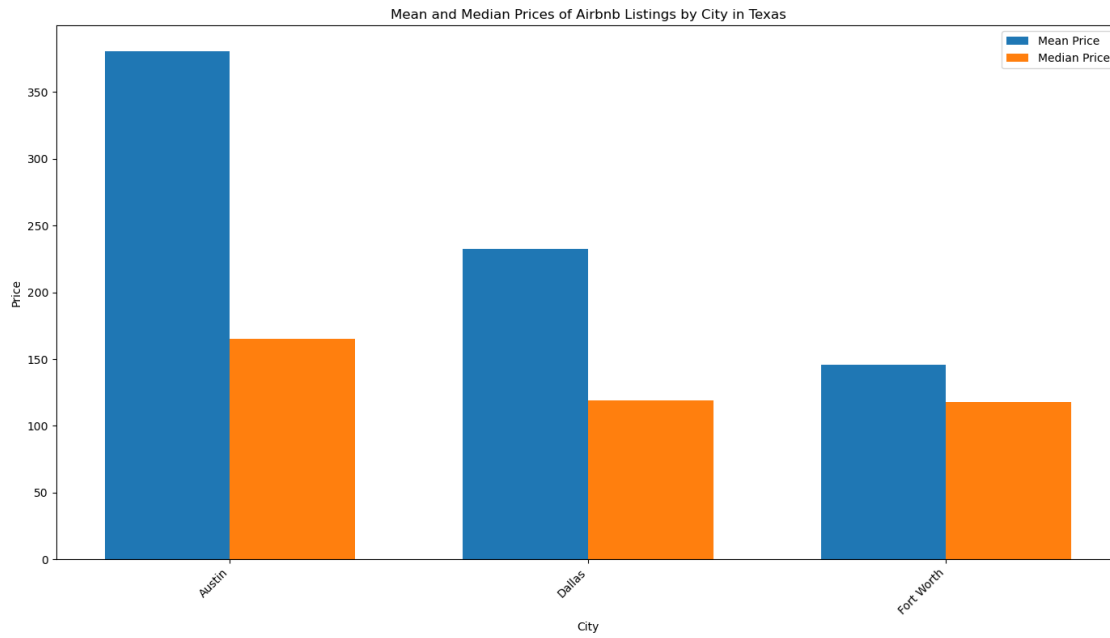
# Define the positions for the bars
x = np.arange(len(california_cities['city']))
width = 0.35

# Create the plot
plt.figure(figsize=(14, 8))
plt.bar(x - width/2, california_cities['mean'], width, label='Mean Price')
plt.bar(x + width/2, california_cities['median'], width, label='Median Price')

# Add labels, title, and legend
plt.xlabel('City')
plt.ylabel('Price')
plt.title('Mean and Median Prices of Airbnb Listings by City in Texas')
plt.xticks(ticks=x, labels=california_cities['city'], rotation=45, ha='right')
plt.legend()

# Show the plot
plt.tight_layout()
plt.show()

```



2.1.2 Visualizations

```
[16]: import folium
      from folium.plugins import MarkerCluster
```



```

# Filter the DataFrame for listings in Texas
texas = df[df['state'] == 'TX']

# Create a base map centered around DFW (Dallas-Fort Worth)
m = folium.Map(location=[32.7079, -96.9209], zoom_start=9)

# Create a marker cluster
marker_cluster = MarkerCluster().add_to(m)

# Add markers to the map
for idx, row in texas.iterrows():
    folium.Marker(
        location=[row['latitude'], row['longitude']],
        popup=folium.Popup(
            f"<strong>Price:</strong> ${row['price']}<br>"
            f"<strong>Name:</strong> {row['name']}<br>"
            f"<strong>Min Nights:</strong> {row['minimum_nights']}",
            max_width=300
        ),
        icon=folium.Icon(color='blue', icon='info-sign')
    ).add_to(marker_cluster)

# Add zip code boundaries
# Replace 'path_to_zip_codes.geojson' with the path to your GeoJSON file
folium.GeoJson(
    "texas-zip-codes-_1613.geojson",
    name='Zip Codes',
    style_function=lambda x: {'fillColor': 'transparent', 'color': 'black',
        ↪ 'weight': 1}
).add_to(m)

# Add layer control to toggle GeoJson overlay
folium.LayerControl().add_to(m)

m

```

[16]: <folium.folium.Map at 0x142d5de10>

The interactive map above allows you to look across the map at all Airbnb listings, when you click on a listing it will display the name, price per night, and minimum nights required to book. This gives us a great idea of the competition in certain areas. For someone wanting to invest in an Airbnb property it would be great to look at competition in the area and see also what commonalities the competition has in terms of price, minimum nights.

```

[17]: # Ensure we are working with a fresh copy of the Texas data
texas = df[df['state'] == 'TX'].copy()

# Group by city and calculate mean, median, and count of prices

```

```

texas_cities = texas.groupby('city')['price'].agg(['mean', 'median', 'count']).
    ↪reset_index()

# Define the positions for the bars
x = np.arange(len(texas_cities['city']))
width = 0.35

# Calculate median listing count for Texas
median_listing_count_tx = texas_cities['count'].median()

# Create the plots
plt.figure(figsize=(18, 8))

# Plot mean and median prices for Texas
plt.subplot(1, 2, 1)
bars1 = plt.bar(x - width/2, texas_cities['mean'], width, label='Mean Price')
bars2 = plt.bar(x + width/2, texas_cities['median'], width, label='Median_
    ↪Price')

# Highlight the median price with a horizontal line
median_price_value_tx = texas_cities['median'].median()
plt.axhline(median_price_value_tx, color='red', linestyle='--', label=f'Median_
    ↪Price Line: ${median_price_value_tx:.2f}')

plt.xlabel('City')
plt.ylabel('Price')
plt.title('Mean and Median Prices of Airbnb Listings by City in Texas')
plt.xticks(ticks=x, labels=texas_cities['city'], rotation=45, ha='right')
plt.legend()

# Plot count of listings for Texas
plt.subplot(1, 2, 2)
plt.bar(x, texas_cities['count'], width, color='skyblue')
plt.xlabel('City')
plt.ylabel('Number of Listings')
plt.title('Number of Airbnb Listings by City in Texas')
plt.xticks(ticks=x, labels=texas_cities['city'], rotation=45, ha='right')

# Add a horizontal line for the median number of listings
plt.axhline(median_listing_count_tx, color='blue', linestyle='--',
    ↪label=f'Median Listing Count Line: {median_listing_count_tx:.0f}')
plt.legend()

plt.tight_layout()
plt.show()

# Ensure we are working with a fresh copy of the California data

```

```

california = df[df['state'] == 'CA'].copy()

# Group by city and calculate mean, median, and count of prices
california_cities = california.groupby('city')['price'].agg(['mean', 'median', 'count']).reset_index()

# Define the positions for the bars
x = np.arange(len(california_cities['city']))
width = 0.35

# Calculate median listing count for California
median_listing_count_ca = california_cities['count'].median()

# Create the plots
plt.figure(figsize=(18, 8))

# Plot mean and median prices for California
plt.subplot(1, 2, 1)
bars1 = plt.bar(x - width/2, california_cities['mean'], width, label='Mean Price')
bars2 = plt.bar(x + width/2, california_cities['median'], width, label='Median Price')

# Highlight the median price with a horizontal line
median_price_value_ca = california_cities['median'].median()
plt.axhline(median_price_value_ca, color='red', linestyle='--', label=f'Median Price Line: ${median_price_value_ca:.2f}')

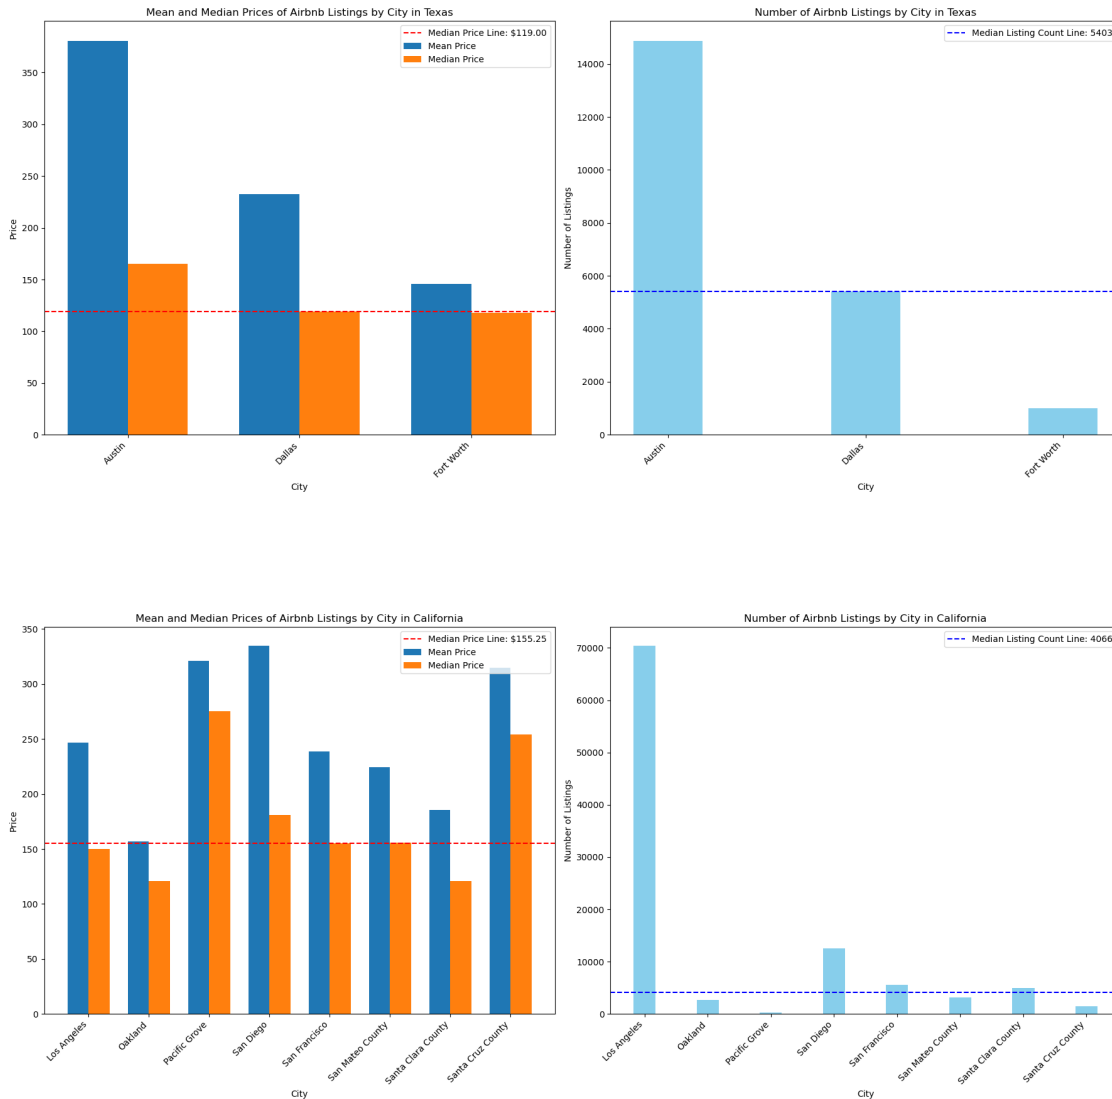
plt.xlabel('City')
plt.ylabel('Price')
plt.title('Mean and Median Prices of Airbnb Listings by City in California')
plt.xticks(ticks=x, labels=california_cities['city'], rotation=45, ha='right')
plt.legend()

# Plot count of listings for California
plt.subplot(1, 2, 2)
plt.bar(x, california_cities['count'], width, color='skyblue')
plt.xlabel('City')
plt.ylabel('Number of Listings')
plt.title('Number of Airbnb Listings by City in California')
plt.xticks(ticks=x, labels=california_cities['city'], rotation=45, ha='right')

# Add a horizontal line for the median number of listings
plt.axhline(median_listing_count_ca, color='blue', linestyle='--', label=f'Median Listing Count Line: {median_listing_count_ca:.0f}')
plt.legend()

```

```
plt.tight_layout()
plt.show()
```



2.1.3 Findings on the Most Popular Airbnb Properties in Texas

Based on our analysis of the Airbnb dataset, we created a popularity score for properties in Texas by considering the number of reviews, reviews per month, and availability. This score was calculated as follows:

$$\text{popularity_score} = \text{number_of_reviews} + (\text{reviews_per_month} * 30) - (\text{availability_365} / 10)$$

After sorting the properties by their popularity scores, we identified the top 10 most popular Airbnb properties in Texas. Here are the key findings:

1. **Top Properties:** The most popular properties are located in various cities across Texas,

highlighting the widespread interest in Airbnb accommodations throughout the state.

2. **High Review Count:** These properties have a high number of reviews, indicating that they are frequently booked and reviewed by guests.
3. **Active Engagement:** Properties with higher reviews per month are consistently attracting guests, suggesting active engagement and positive experiences that encourage reviews.
4. **Low Availability:** Many of the top properties have lower availability throughout the year, implying they are frequently booked and in high demand.

This analysis provides valuable insights into the factors contributing to the popularity of Airbnb properties. Hosts and property managers can leverage this information to enhance their listings, improve guest experiences, and increase bookings by focusing on generating positive reviews and maintaining high occupancy rates.

```
[18]: # Ensure we are working with a fresh copy
texas = df[df['state'] == 'TX'].copy()

# Create a popularity score using .loc
texas.loc[:, 'popularity_score'] = texas['number_of_reviews'] +_
    ↪(texas['reviews_per_month'] * 30) - (texas['availability_365'] / 10)

# Sort properties by popularity score
most_popular_properties_texas = texas.sort_values(by='popularity_score',_
    ↪ascending=False).reset_index()

# Select the top 10 most popular properties
top_10_most_popular_texas = most_popular_properties_texas.head(10)

# Define colors for different cities in Texas
city_colors_texas = {
    'Austin': 'lightskyblue',
    'Dallas': 'lightgoldenrodyellow',
    'Houston': 'lightgreen',
    # Add more cities if needed
}


# Plot the top 10 most popular properties in Texas
plt.figure(figsize=(12, 6))
colors_texas = top_10_most_popular_texas['city'].map(city_colors_texas).
    ↪fillna('grey') # Handle missing colors
bars_texas = plt.barh(top_10_most_popular_texas['name'],_
    ↪top_10_most_popular_texas['popularity_score'], color=colors_texas)

# Create a legend
handles_texas = [plt.Line2D([0], [0], color=color, lw=4) for color in_
    ↪city_colors_texas.values()]
labels_texas = city_colors_texas.keys()
plt.legend(handles_texas, labels_texas, title='City', loc='best')
```

```
plt.xlabel('Popularity Score')
plt.title('Top 10 Most Popular Airbnb Properties in Texas')
plt.gca().invert_yaxis()
plt.show()
```

/Users/brianhonea/anaconda3/lib/python3.11/site-packages/IPython/core/pylabtools.py:152: UserWarning: Glyph 10024 (\N{SPARKLES}) missing from current font.

```
fig.canvas.print_figure(bytes_io, **kw)
```



```
[19]: # Ensure we are working with a fresh copy
california = df[df['state'] == 'CA'].copy()

# Create a popularity score using .loc
california.loc[:, 'popularity_score'] = california['number_of_reviews'] +
    (california['reviews_per_month'] * 30) - (california['availability_365'] /
    10)

# Sort properties by popularity score
most_popular_properties_california = california.
    sort_values(by='popularity_score', ascending=False).reset_index()

# Select the top 10 most popular properties
top_10_most_popular_california = most_popular_properties_california.head(15)

# Define colors for different cities in California
city_colors_california = {
    'Los Angeles': 'lightpink',
    'San Francisco': 'lightcoral',
    'Santa Cruz County': 'lightgreen',
    'Pacific Grove': 'lightsalmon',
    # Add more cities if needed
}
```

```

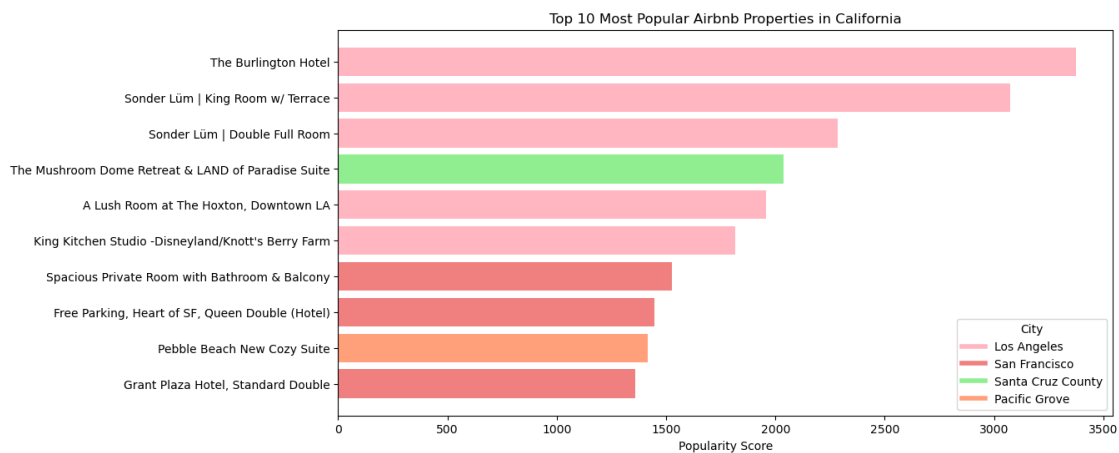
}

# Plot the top 10 most popular properties in California
plt.figure(figsize=(12, 6))
colors_california = top_10_most_popular_california['city'].
    ↳map(city_colors_california).fillna('grey') # Handle missing colors
bars_california = plt.barh(top_10_most_popular_california['name'],
    ↳top_10_most_popular_california['popularity_score'], color=colors_california)

# Create a legend
handles_california = [plt.Line2D([0], [0], color=color, lw=4) for color in
    ↳city_colors_california.values()]
labels_california = city_colors_california.keys()
plt.legend(handles_california, labels_california, title='City', loc='best')

plt.xlabel('Popularity Score')
plt.title('Top 10 Most Popular Airbnb Properties in California')
plt.gca().invert_yaxis()
plt.show()

```



```

[20]: # Calculate the count of listings per neighborhood
neighbourhood_counts = california['neighbourhood'].value_counts().reset_index()
neighbourhood_counts.columns = ['neighbourhood', 'count']

# Select the top 10 neighborhoods by count
top_10_neighbourhoods = neighbourhood_counts.head(10)

# Filter the original DataFrame to include only the top 10 neighborhoods
top_10_data = california[california['neighbourhood'].
    ↳isin(top_10_neighbourhoods['neighbourhood'])]

```

```

# Group by neighborhood and calculate mean and median prices
top_10_prices = top_10_data.groupby('neighbourhood')['price'].agg(['mean',
↳ 'median']).reset_index()

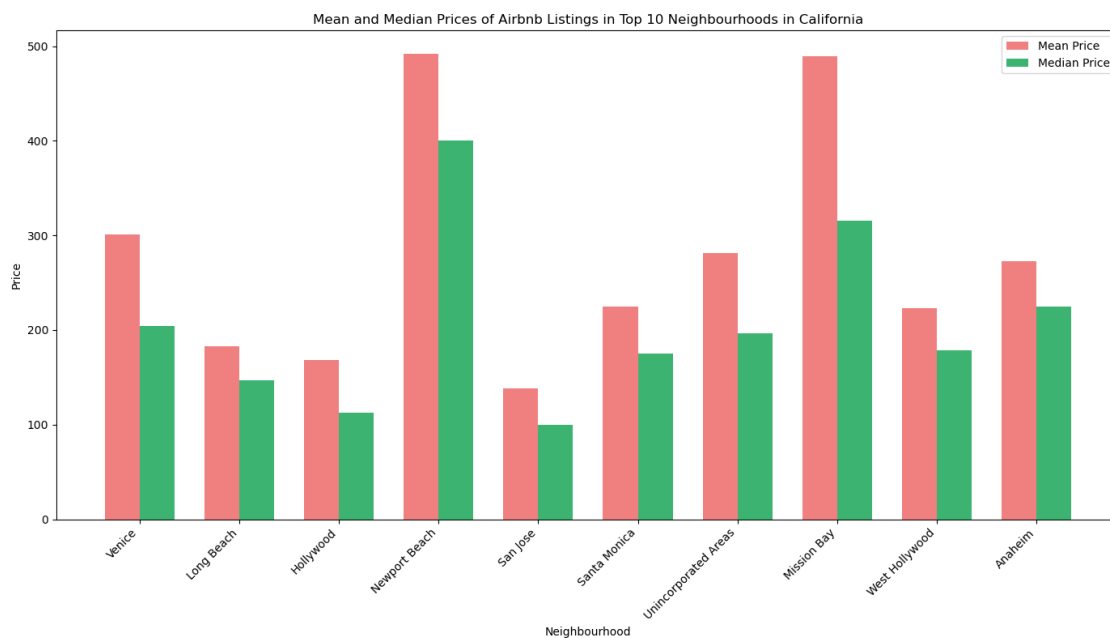
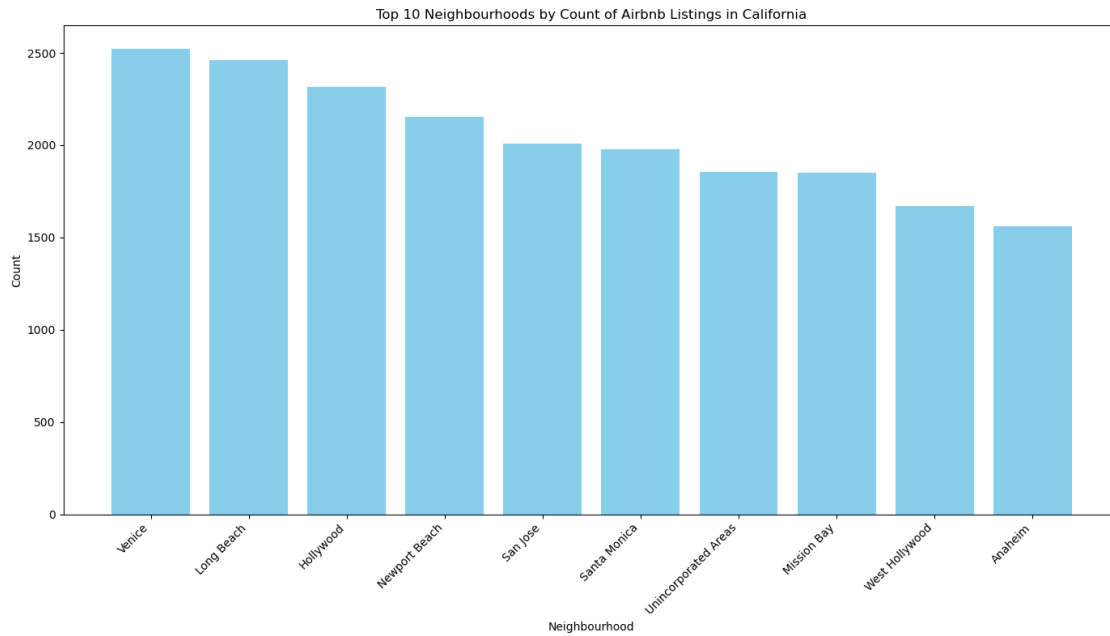
# Sort the mean and median prices DataFrame to match the order of the top 10
↳ neighborhoods by count
top_10_prices = top_10_prices.set_index('neighbourhood').
↳ loc[top_10_neighbourhoods['neighbourhood']].reset_index()

# Create the first plot for the top 10 neighborhoods by count
plt.figure(figsize=(14, 8))
plt.bar(top_10_neighbourhoods['neighbourhood'], top_10_neighbourhoods['count'],
↳ color='skyblue')
plt.xlabel('Neighbourhood')
plt.ylabel('Count')
plt.title('Top 10 Neighbourhoods by Count of Airbnb Listings in California')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()

# Create the second plot for the mean and median prices of the top 10
↳ neighborhoods
x = np.arange(len(top_10_prices['neighbourhood']))
width = 0.35

plt.figure(figsize=(14, 8))
plt.bar(x - width/2, top_10_prices['mean'], width, label='Mean Price',
↳ color='lightcoral')
plt.bar(x + width/2, top_10_prices['median'], width, label='Median Price',
↳ color='mediumseagreen')
plt.xlabel('Neighbourhood')
plt.ylabel('Price')
plt.title('Mean and Median Prices of Airbnb Listings in Top 10 Neighbourhoods
↳ in California')
plt.xticks(ticks=x, labels=top_10_prices['neighbourhood'], rotation=45,
↳ ha='right')
plt.legend()
plt.tight_layout()
plt.show()

```

```
[21]: # Calculate the count of listings per neighborhood
neighbourhood_counts = texas['neighbourhood'].value_counts().reset_index()
neighbourhood_counts.columns = ['neighbourhood', 'count']

# Select the top 10 neighborhoods by count
```

```

top_10_neighbourhoods = neighbourhood_counts.head(10)

# Filter the original DataFrame to include only the top 10 neighborhoods
top_10_data = texas[texas['neighbourhood']].
    ↪isin(top_10_neighbourhoods['neighbourhood'])]

# Group by neighborhood and calculate mean and median prices
top_10_prices = top_10_data.groupby('neighbourhood')['price'].agg(['mean',
    ↪'median']).reset_index()

# Sort the mean and median prices DataFrame to match the order of the top 10
    ↪neighborhoods by count
top_10_prices = top_10_prices.set_index('neighbourhood').
    ↪loc[top_10_neighbourhoods['neighbourhood']].reset_index()

# Calculate the overall median price for the top 10 neighborhoods
overall_median_price = top_10_data['price'].median()

# Create the first plot for the top 10 neighborhoods by count
plt.figure(figsize=(14, 8))
plt.bar(top_10_neighbourhoods['neighbourhood'], top_10_neighbourhoods['count'],
    ↪color='skyblue')
plt.xlabel('Neighbourhood')
plt.ylabel('Count')
plt.title('Top 10 Neighbourhoods by Count of Airbnb Listings in Texas')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()

# Create the second plot for the mean and median prices of the top 10
    ↪neighborhoods
x = np.arange(len(top_10_prices['neighbourhood']))
width = 0.35

plt.figure(figsize=(14, 8))
plt.bar(x - width/2, top_10_prices['mean'], width, label='Mean Price',
    ↪color='lightcoral')
plt.bar(x + width/2, top_10_prices['median'], width, label='Median Price',
    ↪color='mediumseagreen')

# Add a horizontal line for the overall median price
plt.axhline(y=overall_median_price, color='blue', linestyle='--', linewidth=1.
    ↪5, label=f'Overall Median Price: ${overall_median_price:.2f}')

# Add labels, title, and legend
plt.xlabel('Neighbourhood')

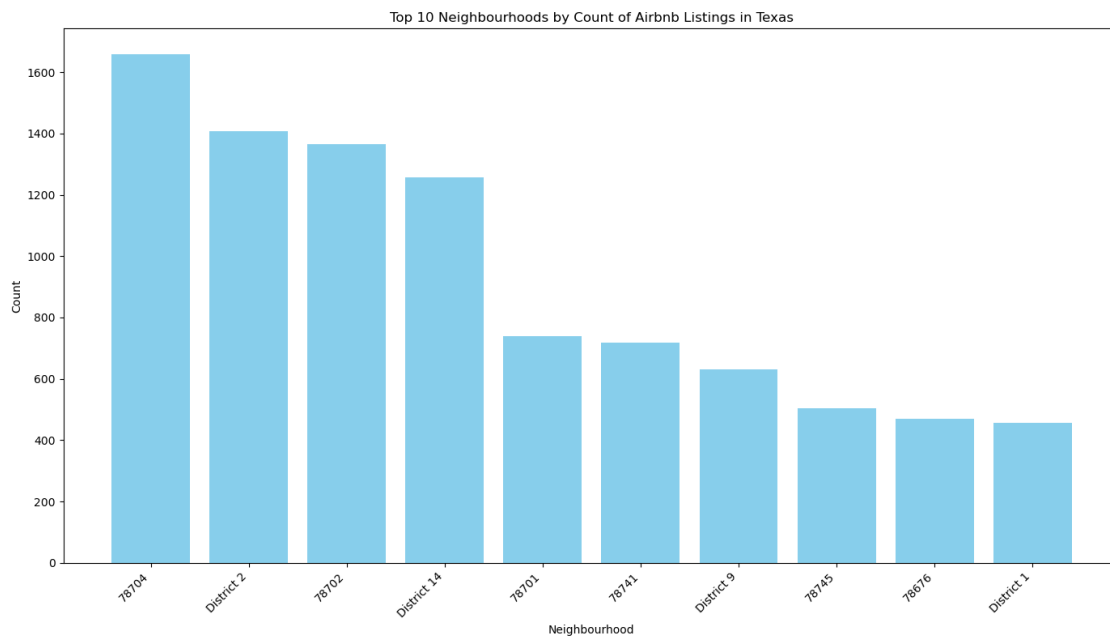
```


```

plt.ylabel('Price')
plt.title('Mean and Median Prices of Airbnb Listings in Top 10 Neighbourhoods_
↳in Texas')
plt.xticks(ticks=x, labels=top_10_prices['neighbourhood'], rotation=45,
↳ha='right')
plt.legend()

# Show the plot
plt.tight_layout()
plt.show()

```





```
[22]: # Filter for Austin
austin_data = texas[texas['city'] == 'Austin']

# Group by neighborhood and calculate mean, median, and count of prices
austin_neighborhoods = austin_data.groupby('neighbourhood')['price'].
    ↪agg(['mean', 'median', 'count']).reset_index()

# Get the top 5 neighborhoods by number of listings
top_5_neighborhoods = austin_neighborhoods.nlargest(5, 'count')

# Define the positions for the bars
x = np.arange(len(top_5_neighborhoods['neighbourhood']))
width = 0.35

# Create the plots
plt.figure(figsize=(18, 8))

# Plot mean and median prices
plt.subplot(1, 2, 1)
bars1 = plt.bar(x - width/2, top_5_neighborhoods['mean'], width, label='Mean_
    ↪Price')
bars2 = plt.bar(x + width/2, top_5_neighborhoods['median'], width,
    ↪label='Median Price')

# Highlight the median price with a horizontal line
median_price_value = top_5_neighborhoods['median'].median()
```

```

plt.axhline(median_price_value, color='red', linestyle='--', label=f'Median
    ↳Price Line: ${median_price_value:.2f}')

plt.xlabel('Neighbourhood')
plt.ylabel('Price')
plt.title('Mean and Median Prices of Airbnb Listings by Neighbourhood in
    ↳Austin')
plt.xticks(ticks=x, labels=top_5_neighborhoods['neighbourhood'], rotation=45,
    ↳ha='right')
plt.legend()


# Plot count of listings
plt.subplot(1, 2, 2)
plt.bar(x, top_5_neighborhoods['count'], width, color='skyblue')
plt.xlabel('Neighborhood')
plt.ylabel('Number of Listings')
plt.title('Number of Airbnb Listings by Neighbourhood in Austin')
plt.xticks(ticks=x, labels=top_5_neighborhoods['neighbourhood'], rotation=45,
    ↳ha='right')

# Highlight the median listing count with a horizontal line
median_listing_count_austin = top_5_neighborhoods['count'].median()

# Add a horizontal line for the median number of listings
plt.axhline(median_listing_count_austin, color='blue', linestyle='--',
    ↳label=f'Median Listing Count Line: {median_listing_count_austin:.0f}')
plt.legend()

plt.tight_layout()
plt.show()

```



```

[23]: # Filter for Dallas
dallas_data = texas[texas['city'] == 'Dallas']

# Group by neighbourhood and calculate mean, median, and count of prices
dallas_neighbourhoods = dallas_data.groupby('neighbourhood')['price'].
    .agg(['mean', 'median', 'count']).reset_index()

# Get the top 5 neighbourhoods by number of listings
top_5_neighbourhoods_dallas = dallas_neighbourhoods.nlargest(5, 'count')

# Define the positions for the bars
x = np.arange(len(top_5_neighbourhoods_dallas['neighbourhood']))
width = 0.35

# Create the plots
plt.figure(figsize=(18, 8))

# Plot mean and median prices
plt.subplot(1, 2, 1)
bars1 = plt.bar(x - width/2, top_5_neighbourhoods_dallas['mean'], width,
    label='Mean Price')
bars2 = plt.bar(x + width/2, top_5_neighbourhoods_dallas['median'], width,
    label='Median Price')

# Highlight the median price with a horizontal line
median_price_value_dallas = top_5_neighbourhoods_dallas['median'].median()
plt.axhline(median_price_value_dallas, color='red', linestyle='--',
    label=f'Median Price Line: ${median_price_value_dallas:.2f}')

plt.xlabel('Neighbourhood')
plt.ylabel('Price')
plt.title('Mean and Median Prices of Airbnb Listings by Neighbourhood in
    Dallas')
plt.xticks(ticks=x, labels=top_5_neighbourhoods_dallas['neighbourhood'],
    rotation=45, ha='right')
plt.legend()

# Plot count of listings
plt.subplot(1, 2, 2)
plt.bar(x, top_5_neighbourhoods_dallas['count'], width, color='skyblue')
plt.xlabel('Neighbourhood')
plt.ylabel('Number of Listings')
plt.title('Number of Airbnb Listings by Neighbourhood in Dallas')


```

```
plt.xticks(ticks=x, labels=top_5_neighbourhoods_dallas['neighbourhood'],
           rotation=45, ha='right')

# Highlight the median listing count with a horizontal line
median_listing_count_dallas = top_5_neighbourhoods_dallas['count'].median()

# Add a horizontal line for the median number of listings
plt.axhline(median_listing_count_dallas, color='blue', linestyle='--',
            label=f'Median Listing Count Line: {median_listing_count_dallas:.0f}')
plt.legend()

plt.tight_layout()
plt.show()
```



```
[24]: # Filter for Fort Worth
fort_worth_data = texas[texas['city'] == 'Fort Worth']

# Group by neighbourhood and calculate mean, median, and count of prices
fort_worth_neighbourhoods = fort_worth_data.groupby('neighbourhood')['price'].
    .agg(['mean', 'median', 'count']).reset_index()

# Get the top 5 neighbourhoods by number of listings
top_5_neighbourhoods_fort_worth = fort_worth_neighbourhoods.nlargest(5, 'count')

# Define the positions for the bars
x = np.arange(len(top_5_neighbourhoods_fort_worth['neighbourhood']))
width = 0.35

# Create the plots
plt.figure(figsize=(18, 8))
```

```

# Plot mean and median prices
plt.subplot(1, 2, 1)
bars1 = plt.bar(x - width/2, top_5_neighbourhoods_fort_worth['mean'], width,
    ↪label='Mean Price')
bars2 = plt.bar(x + width/2, top_5_neighbourhoods_fort_worth['median'], width,
    ↪label='Median Price')

# Highlight the median price with a horizontal line
median_price_value_fort_worth = top_5_neighbourhoods_fort_worth['median'].
    ↪median()
plt.axhline(median_price_value_fort_worth, color='red', linestyle='--',
    ↪label=f'Median Price Line: ${median_price_value_fort_worth:.2f}')

plt.xlabel('Neighbourhood')
plt.ylabel('Price')
plt.title('Mean and Median Prices of Airbnb Listings by Neighbourhood in Fort
    ↪Worth')
plt.xticks(ticks=x, labels=top_5_neighbourhoods_fort_worth['neighbourhood'],
    ↪rotation=45, ha='right')
plt.legend()


# Plot count of listings
plt.subplot(1, 2, 2)
plt.bar(x, top_5_neighbourhoods_fort_worth['count'], width, color='skyblue')
plt.xlabel('Neighbourhood')
plt.ylabel('Number of Listings')
plt.title('Number of Airbnb Listings by Neighbourhood in Fort Worth')
plt.xticks(ticks=x, labels=top_5_neighbourhoods_fort_worth['neighbourhood'],
    ↪rotation=45, ha='right')

# Highlight the median listing count with a horizontal line
median_listing_count_fort_worth = top_5_neighbourhoods_fort_worth['count'].
    ↪median()

# Add a horizontal line for the median number of listings
plt.axhline(median_listing_count_fort_worth, color='blue', linestyle='--',
    ↪label=f'Median Listing Count Line: {median_listing_count_fort_worth:.0f}')
plt.legend()

plt.tight_layout()
plt.show()

```

2.2 Predictive Modeling

At this point in the analysis we have leveraged our visual and honed in on Texas as a viable investment market. Specifically we have chosen the following:

Dallas and Fort Worth

```
[25]: # Calculate average price per city
texas['avg_price_city'] = texas.groupby('city')['price'].transform('mean')

# Calculate average number of reviews per city
texas['avg_reviews_city'] = texas.groupby('city')['number_of_reviews'].
    ↪transform('mean')

# Calculate availability score
texas['availability_score'] = texas['availability_365'] / 365

# Create feature matrix and target vector
features = ['latitude', 'longitude', 'room_type', 'minimum_nights',
            'avg_price_city', 'avg_reviews_city', 'availability_score', 'city']
target = 'price'

X = texas[features]
y = texas[target]
```

```
[26]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score
```

```

# Define preprocessing pipelines
numeric_features = ['latitude', 'longitude', 'minimum_nights',
                    'avg_price_city', 'avg_reviews_city', 'availability_score']
categorical_features = ['room_type', 'city']

numeric_pipeline = Pipeline(steps=[
    ('scaler', StandardScaler())
])

categorical_pipeline = Pipeline(steps=[
    ('onehot', OneHotEncoder(handle_unknown='ignore'))
])

preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_pipeline, numeric_features),
        ('cat', categorical_pipeline, categorical_features)
    ]
)

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
                                                    random_state=42)

# Build and train the model
model_pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('regressor', RandomForestRegressor(random_state=42))
])

model_pipeline.fit(X_train, y_train)

# Make predictions
y_pred = model_pipeline.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f'Mean Squared Error: {mse}')
print(f'R^2 Score: {r2}')

```

Mean Squared Error: 163683.01537288324
R^2 Score: 0.8692285285926519

```
[27]: # Add predictions to the dataset
texas['predicted_price'] = model_pipeline.predict(X)


# Calculate average predicted price per city
avg_predicted_price_city = texas.groupby('city')['predicted_price'].mean().
    ↪reset_index()
avg_predicted_price_city = avg_predicted_price_city.
    ↪sort_values(by='predicted_price', ascending=False)

# Get top 10 cities
top_cities = avg_predicted_price_city.head(10)

print("Top 10 cities for investment:")
print(top_cities)
```

```
Top 10 cities for investment:
      city  predicted_price
0   Austin      386.279324
1   Dallas      237.854439
2  Fort Worth      155.381815
```

```
[28]: # Plot top 10 cities for investment
plt.figure(figsize=(12, 8))
sns.barplot(x='predicted_price', y='city', data=top_cities)
plt.title('Top 10 Cities for Investment')
plt.xlabel('Predicted Average Rental Price')
plt.ylabel('City')
plt.show()
```



```
[29]: #Feature Engineering

# Calculate average price per neighborhood
texas['avg_price_neigh'] = texas.groupby('neighbourhood')['price'].
    ↪transform('mean')

# Calculate average number of reviews per neighborhood
texas['avg_reviews_neigh'] = texas.
    ↪groupby('neighbourhood')['number_of_reviews'].transform('mean')

# Calculate availability score
texas['availability_score'] = texas['availability_365'] / 365

# Create feature matrix and target vector
features = ['latitude', 'longitude', 'room_type', 'minimum_nights',
            'avg_price_neigh', 'avg_reviews_neigh', 'availability_score',
            ↪'city']
target = 'price'

X = texas[features]
y = texas[target]
```

```
[30]: # Calculate average price per city
texas['avg_price_city'] = texas.groupby('city')['price'].transform('mean')

# Calculate average number of reviews per city
texas['avg_reviews_city'] = texas.groupby('city')['number_of_reviews'].
    ↪transform('mean')

# Calculate availability score
texas['availability_score'] = texas['availability_365'] / 365

# Create feature matrix and target vector
features = ['latitude', 'longitude', 'room_type', 'minimum_nights',
            'avg_price_city', 'avg_reviews_city', 'availability_score', 'city']
target = 'price'

X = texas[features]
y = texas[target]
```

```
[31]: # Add predictions to the dataset
texas['predicted_price'] = model_pipeline.predict(X)

# Calculate average predicted price per city
avg_predicted_price_city = texas.groupby('city')['predicted_price'].mean().
    ↪reset_index()
avg_predicted_price_city = avg_predicted_price_city.
    ↪sort_values(by='predicted_price', ascending=False)

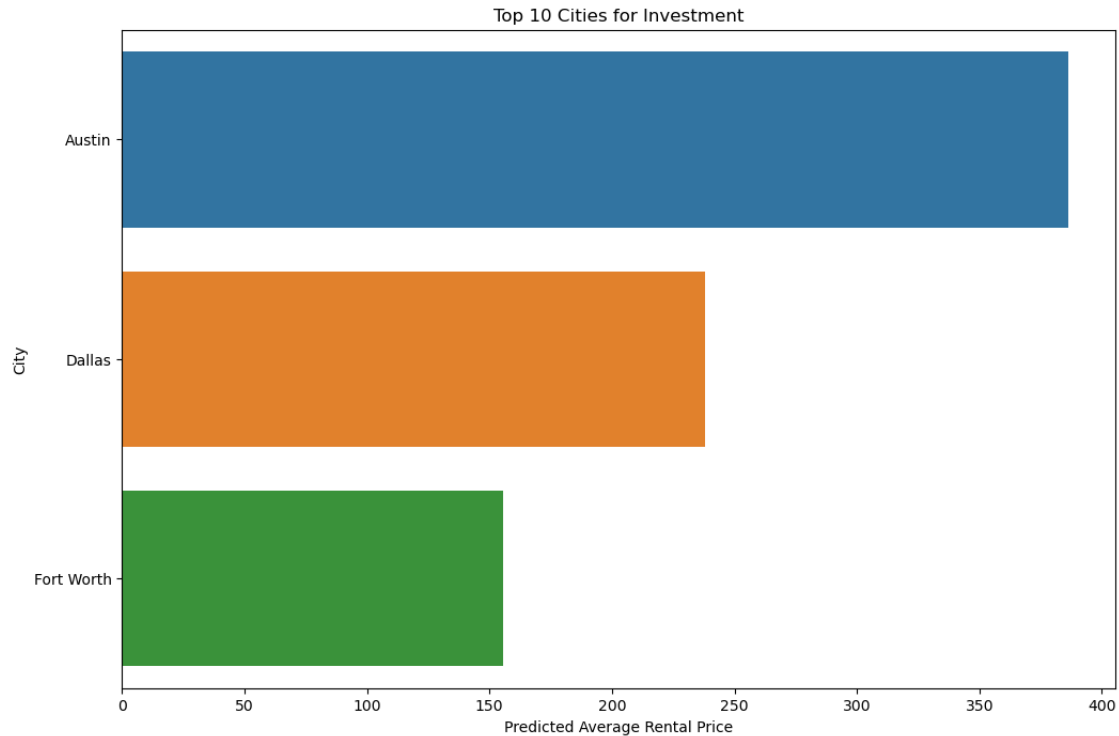
# Get top 10 cities
top_cities = avg_predicted_price_city.head(10)

print("Top 10 cities for investment:")
print(top_cities)
```

Top 10 cities for investment:

	city	predicted_price
0	Austin	386.279324
1	Dallas	237.854439
2	Fort Worth	155.381815

```
[32]: # Plot top 10 cities for investment
plt.figure(figsize=(12, 8))
sns.barplot(x='predicted_price', y='city', data=top_cities)
plt.title('Top 10 Cities for Investment')
plt.xlabel('Predicted Average Rental Price')
plt.ylabel('City')
plt.show()
```



```
[33]: # Assuming top_cities DataFrame is available from the previous steps
top_city = top_cities['city'].iloc[0] # Get the top city
df_city = texas[texas['city'] == top_city]

# Calculate average price per neighborhood
df_city['avg_price_neigh'] = df_city.groupby('neighbourhood')['price'].
    ↪transform('mean')

# Calculate average number of reviews per neighborhood
df_city['avg_reviews_neigh'] = df_city.
    ↪groupby('neighbourhood')['number_of_reviews'].transform('mean')

# Calculate availability score
df_city['availability_score'] = df_city['availability_365'] / 365

# Create feature matrix and target vector
features_neigh = ['latitude', 'longitude', 'room_type', 'minimum_nights',
                  'avg_price_neigh', 'avg_reviews_neigh', 'availability_score']
target_neigh = 'price'

X_neigh = df_city[features_neigh]
y_neigh = df_city[target_neigh]
```

```

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score

# Define preprocessing pipelines
numeric_features_neigh = ['latitude', 'longitude', 'minimum_nights',
                           'avg_price_neigh', 'avg_reviews_neigh',
                           'availability_score']
categorical_features_neigh = ['room_type']

numeric_pipeline_neigh = Pipeline(steps=[
    ('scaler', StandardScaler())
])

categorical_pipeline_neigh = Pipeline(steps=[
    ('onehot', OneHotEncoder(handle_unknown='ignore'))
])

preprocessor_neigh = ColumnTransformer(
    transformers=[
        ('num', numeric_pipeline_neigh, numeric_features_neigh),
        ('cat', categorical_pipeline_neigh, categorical_features_neigh)
    ]
)

# Train-test split
X_train_neigh, X_test_neigh, y_train_neigh, y_test_neigh =
    train_test_split(X_neigh, y_neigh, test_size=0.3, random_state=42)

# Build and train the model
model_pipeline_neigh = Pipeline(steps=[
    ('preprocessor', preprocessor_neigh),
    ('regressor', RandomForestRegressor(random_state=42))
])

model_pipeline_neigh.fit(X_train_neigh, y_train_neigh)

# Make predictions
y_pred_neigh = model_pipeline_neigh.predict(X_test_neigh)

# Evaluate the model
mse_neigh = mean_squared_error(y_test_neigh, y_pred_neigh)
r2_neigh = r2_score(y_test_neigh, y_pred_neigh)

```

```

print(f'Mean Squared Error: {mse_neigh}')
print(f'R^2 Score: {r2_neigh}')

# Add predictions to the dataset
df_city['predicted_price'] = model_pipeline_neigh.predict(X_neigh)

# Calculate average predicted price per neighborhood
avg_predicted_price_neigh = df_city.groupby('neighbourhood')['predicted_price'].
    ↪mean().reset_index()
avg_predicted_price_neigh = avg_predicted_price_neigh.
    ↪sort_values(by='predicted_price', ascending=False)

# Get top 10 neighborhoods
top_neigh = avg_predicted_price_neigh.head(10)

print(f"Top 10 neighborhoods for investment in {top_city}:")
print(top_neigh)

# Plot top 10 neighborhoods for investment in the selected city
plt.figure(figsize=(12, 8))
sns.barplot(x='predicted_price', y='neighbourhood', data=top_neigh)
plt.title(f'Top 10 Neighborhoods for Investment in {top_city}')
plt.xlabel('Predicted Average Rental Price')
plt.ylabel('Neighborhood')
plt.show()

```

/var/folders/rm/wtgg2sx94jq9c87_mpvkj700000gn/T/ipykernel_27801/62496316.py:6:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_city['avg_price_neigh'] =
df_city.groupby('neighbourhood')['price'].transform('mean')
```

/var/folders/rm/wtgg2sx94jq9c87_mpvkj700000gn/T/ipykernel_27801/62496316.py:9:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_city['avg_reviews_neigh'] =
df_city.groupby('neighbourhood')['number_of_reviews'].transform('mean')
```

/var/folders/rm/wtgg2sx94jq9c87_mpvkj700000gn/T/ipykernel_27801/62496316.py:12:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_city['availability_score'] = df_city['availability_365'] / 365
```

Mean Squared Error: 141546.71300135995

R² Score: 0.8973139425732697

Top 10 neighborhoods for investment in Austin:

	neighbourhood	predicted_price
70	78750	1924.693800
57	78735	1058.856417
78	78759	1024.308203
66	78746	1015.013500
49	78727	931.417881
54	78732	919.099000
52	78730	797.261154
53	78731	788.439840
77	78758	763.950419
62	78741	679.123314

/var/folders/rm/wtgg2sx94jq9c87_mpvkj700000gn/T/ipykernel_27801/62496316.py:72:

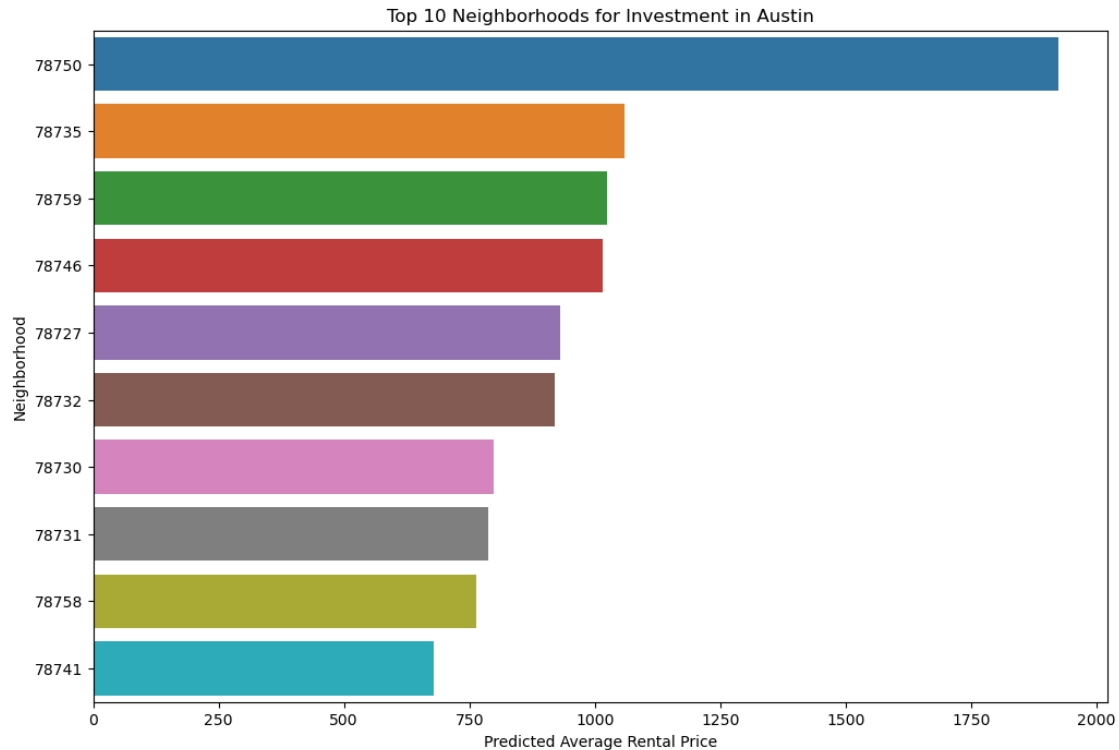
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_city['predicted_price'] = model_pipeline_neigh.predict(X_neigh)
```



```
[34]: #Do the same for dallas and fort worth

# Filter for state = TX and specific cities
texas = df[df['state'] == 'TX']
df_dallas = texas[texas['city'] == 'Dallas']
df_fortworth = texas[texas['city'] == 'Fort Worth']

# For Dallas
df_dallas['avg_price_neigh'] = df_dallas.groupby('neighbourhood')['price'].
    ↪transform('mean')
df_dallas['avg_reviews_neigh'] = df_dallas.
    ↪groupby('neighbourhood')['number_of_reviews'].transform('mean')
df_dallas['availability_score'] = df_dallas['availability_365'] / 365

# For Fort Worth
df_fortworth['avg_price_neigh'] = df_fortworth.
    ↪groupby('neighbourhood')['price'].transform('mean')
df_fortworth['avg_reviews_neigh'] = df_fortworth.
    ↪groupby('neighbourhood')['number_of_reviews'].transform('mean')
df_fortworth['availability_score'] = df_fortworth['availability_365'] / 365
```

```

# Features and target for Dallas
features_neigh = ['latitude', 'longitude', 'room_type', 'minimum_nights',
                  'avg_price_neigh', 'avg_reviews_neigh', 'availability_score']
target_neigh = 'price'

X_dallas = df_dallas[features_neigh]
y_dallas = df_dallas[target_neigh]

# Features and target for Fort Worth
X_fortworth = df_fortworth[features_neigh]
y_fortworth = df_fortworth[target_neigh]

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score

# Define preprocessing pipelines
numeric_features_neigh = ['latitude', 'longitude', 'minimum_nights',
                          'avg_price_neigh', 'avg_reviews_neigh',
                          'availability_score']
categorical_features_neigh = ['room_type']

numeric_pipeline_neigh = Pipeline(steps=[
    ('scaler', StandardScaler())
])

categorical_pipeline_neigh = Pipeline(steps=[
    ('onehot', OneHotEncoder(handle_unknown='ignore'))
])

preprocessor_neigh = ColumnTransformer(
    transformers=[
        ('num', numeric_pipeline_neigh, numeric_features_neigh),
        ('cat', categorical_pipeline_neigh, categorical_features_neigh)
    ]
)

# Model training function
def train_model(X, y):
    # Train-test split
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
    random_state=42)

    # Build and train the model

```

```

model_pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor_neigh),
    ('regressor', RandomForestRegressor(random_state=42))
])

model_pipeline.fit(X_train, y_train)

# Make predictions
y_pred = model_pipeline.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f'Mean Squared Error: {mse}')
print(f'R^2 Score: {r2}')

return model_pipeline

# Train models for Dallas and Fort Worth
model_dallas = train_model(X_dallas, y_dallas)
model_fortworth = train_model(X_fortworth, y_fortworth)

# Add predictions to the datasets
df_dallas['predicted_price'] = model_dallas.predict(X_dallas)
df_fortworth['predicted_price'] = model_fortworth.predict(X_fortworth)

# Calculate average predicted price per neighborhood
avg_predicted_price_neigh_dallas = df_dallas.
    ↳groupby('neighbourhood')['predicted_price'].mean().reset_index()
avg_predicted_price_neigh_dallas = avg_predicted_price_neigh_dallas.
    ↳sort_values(by='predicted_price', ascending=False)

avg_predicted_price_neigh_fortworth = df_fortworth.
    ↳groupby('neighbourhood')['predicted_price'].mean().reset_index()
avg_predicted_price_neigh_fortworth = avg_predicted_price_neigh_fortworth.
    ↳sort_values(by='predicted_price', ascending=False)

# Get top 10 neighborhoods
top_neigh_dallas = avg_predicted_price_neigh_dallas.head(10)
top_neigh_fortworth = avg_predicted_price_neigh_fortworth.head(10)

print("Top 10 neighborhoods for investment in Dallas:")
print(top_neigh_dallas)

print("Top 10 neighborhoods for investment in Fort Worth:")

```

```

print(top_neigh_fortworth)

# Plot top 10 neighborhoods for investment in Dallas
plt.figure(figsize=(12, 8))
sns.barplot(x='predicted_price', y='neighbourhood', data=top_neigh_dallas)
plt.title('Top 10 Neighborhoods for Investment in Dallas')
plt.xlabel('Predicted Average Rental Price')
plt.ylabel('Neighborhood')
plt.show()

# Plot top 10 neighborhoods for investment in Fort Worth
plt.figure(figsize=(12, 8))
sns.barplot(x='predicted_price', y='neighbourhood', data=top_neigh_fortworth)
plt.title('Top 10 Neighborhoods for Investment in Fort Worth')
plt.xlabel('Predicted Average Rental Price')
plt.ylabel('Neighborhood')
plt.show()

```

/var/folders/rm/wtgg2sx94jq9c87_mpvkj700000gn/T/ipykernel_27801/2612041528.py:9
: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df_dallas['avg_price_neigh'] =
df_dallas.groupby('neighbourhood')['price'].transform('mean')
/var/folders/rm/wtgg2sx94jq9c87_mpvkj700000gn/T/ipykernel_27801/2612041528.py:1
0: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df_dallas['avg_reviews_neigh'] =
df_dallas.groupby('neighbourhood')['number_of_reviews'].transform('mean')
/var/folders/rm/wtgg2sx94jq9c87_mpvkj700000gn/T/ipykernel_27801/2612041528.py:1
1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df_dallas['availability_score'] = df_dallas['availability_365'] / 365
/var/folders/rm/wtgg2sx94jq9c87_mpvkj700000gn/T/ipykernel_27801/2612041528.py:1
4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_fortworth['avg_price_neigh'] =  
df_fortworth.groupby('neighbourhood')['price'].transform('mean')  
/var/folders/rm/wtgg2sx94jq9c87_mpvkj700000gn/T/ipykernel_27801/2612041528.py:1  
5: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_fortworth['avg_reviews_neigh'] =  
df_fortworth.groupby('neighbourhood')['number_of_reviews'].transform('mean')  
/var/folders/rm/wtgg2sx94jq9c87_mpvkj700000gn/T/ipykernel_27801/2612041528.py:1  
6: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_fortworth['availability_score'] = df_fortworth['availability_365'] / 365
```

Mean Squared Error: 141619.07048061202

R² Score: 0.7730946576305356

Mean Squared Error: 8465.688019269104

R² Score: 0.21286680295048543

Top 10 neighborhoods for investment in Dallas:

	neighbourhood	predicted_price
12	District 8	337.544800
6	District 2	289.084598
5	District 14	235.023474
3	District 12	216.068317
2	District 11	213.012285
1	District 10	212.544745
4	District 13	202.136601
13	District 9	198.052663
0	District 1	179.120472
10	District 6	171.902486

Top 10 neighborhoods for investment in Fort Worth:

	neighbourhood	predicted_price
1	District 3	174.155172
0	District 2	168.901558
5	District 7	164.387788
2	District 4	159.688976
3	District 5	143.116947
4	District 6	133.353929

```

7    District 9      127.715966
6    District 8      113.474444

```

```

/var/folders/rm/wtgg2sx94jq9c87_mpvkj700000gn/T/ipykernel_27801/2612041528.py:8

```

9: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```

df_dallas['predicted_price'] = model_dallas.predict(X_dallas)

```

```

/var/folders/rm/wtgg2sx94jq9c87_mpvkj700000gn/T/ipykernel_27801/2612041528.py:9

```

0: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

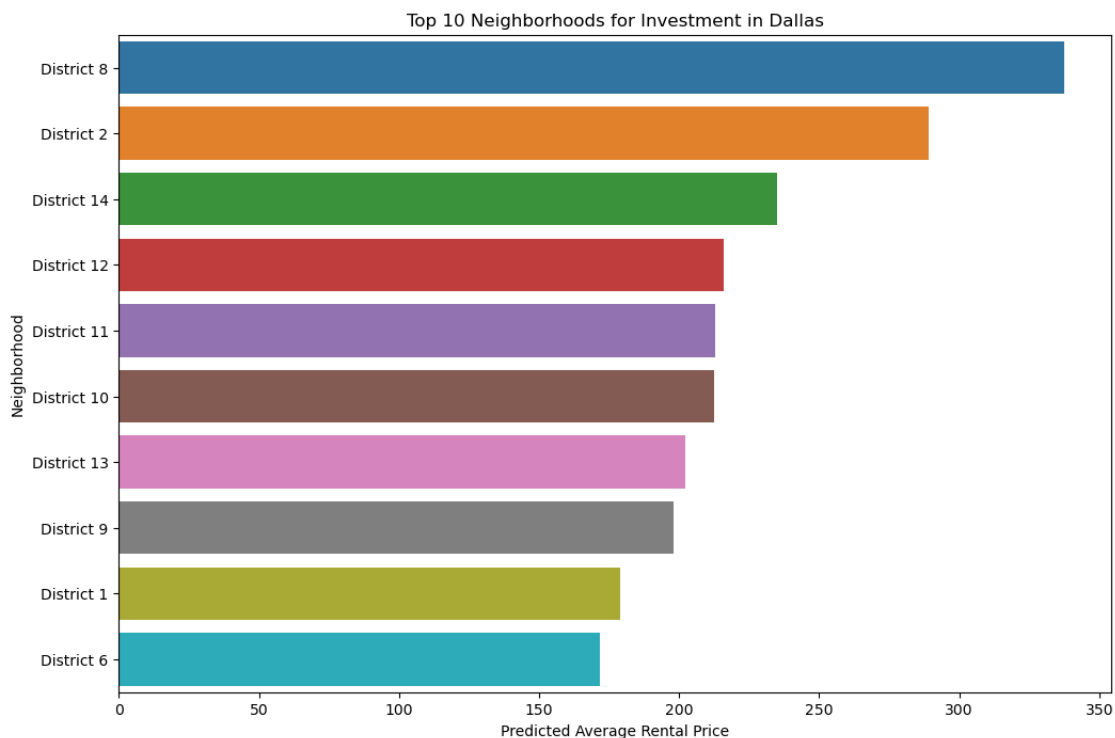
Try using `.loc[row_indexer,col_indexer] = value` instead

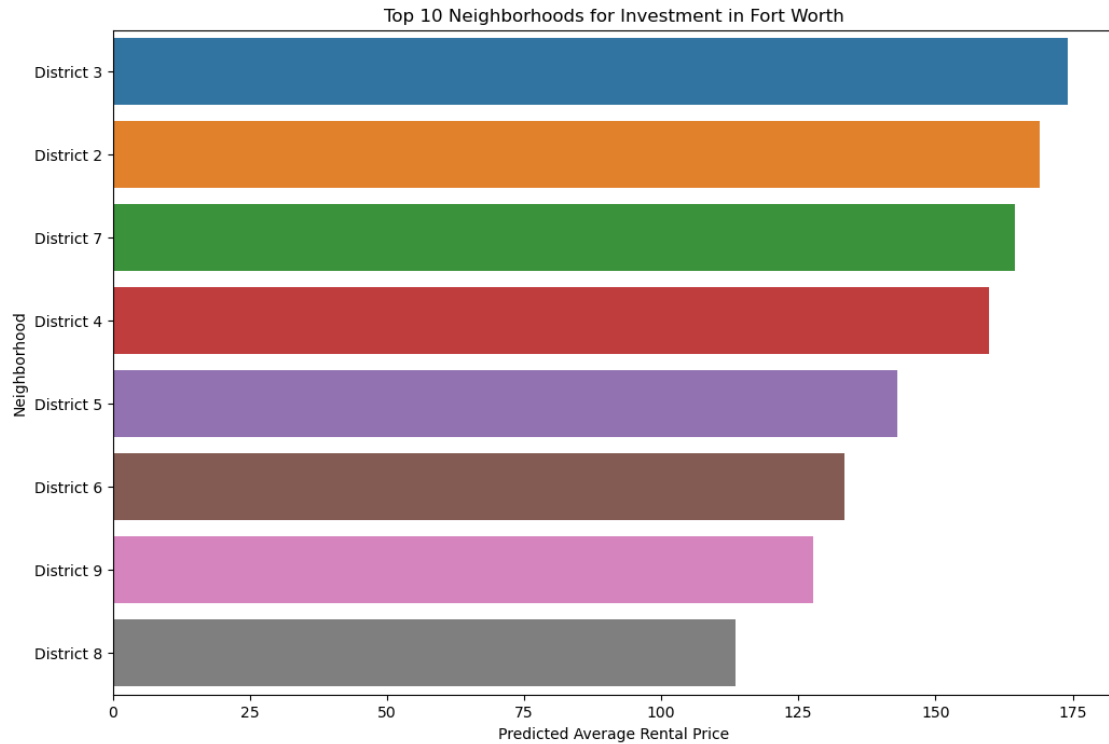
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```

df_fortworth['predicted_price'] = model_fortworth.predict(X_fortworth)

```





Utilizing the predictive modeling, stakeholders will not only be reaffirmed in their decision to invest in Texas Real Estate, but they can also feel additionally confident in picking a neighborhood that poses a substantial return on investment.

This can be done by leveraging the above predicted rental price measures as well as the labeled top neighborhoods to invest in!